

Width 256×6 layers: dropout vs overfitting (same peak LR $1e-3$)

